

Spark SQL Guide

- [Getting Started](#)
- [Data Sources](#)
 - [Generic Load/Save Functions](#)
 - [Parquet Files](#)
 - [ORC Files](#)
 - [JSON Files](#)
 - [Hive Tables](#)
 - [JDBC To Other Databases](#)
 - [Avro Files](#)
 - **[Troubleshooting](#)**
- [Performance Tuning](#)

«

Troubleshooting

- The JDBC driver class must be visible to the primordial class loader on the client session and on all executors. This is because Java's DriverManager class does a security check that results in it ignoring all drivers not visible to the primordial class loader when one goes to open a connection. One convenient way to do this is to modify `compute_classpath.sh` on all worker nodes to include your driver JARs.
- Some databases, such as H2, convert all names to upper case. You'll need to use upper case to refer to those names in Spark SQL.
- Users can specify vendor-specific JDBC connection properties in the data source options to do special treatment. For example, `spark.read.format("jdbc").option("url", oracleJdbcUrl).option("oracle.jdbc.mapDateToTimestamp", "false").oracle.jdbc.mapDateToTimestamp` defaults to `true`, users often need to disable this flag to avoid Oracle date being resolved as timestamp.